

A Modern Biostatistical Re-analysis of the NCCTG Adjuvant Colon Cancer Randomized Trial

Frequentist, Bayesian, and Machine Learning Perspectives

Cathy

September 2026

Abstract

Adjuvant chemotherapy with levamisole and 5-fluorouracil (5-FU) after resection of stage B/C colon carcinoma was established by the North Central Cancer Treatment Group (NCCTG) in the late 1980s, and the resulting individual-patient dataset has become one of the most widely used teaching resources in survival analysis. We ask what a modern biostatistical toolkit can extract from the same trial today. Among 929 patients randomized to observation ($n = 315$), levamisole ($n = 310$), or levamisole plus 5-FU ($n = 304$), we reconstruct a composite disease-free survival (DFS) endpoint (recurrence or death; 506 events) from the published event-type coding and integrate six families of methods in one fully reproducible pipeline: (i) Kaplan–Meier and Cox analyses with multiple imputation, competing-risk sensitivity analyses, and absolute-benefit measures including the restricted mean survival time; (ii) a redesign of the trial as a three-look group-sequential design with O’Brien–Fleming efficacy and non-binding beta-spending futility boundaries, re-enacted on the real event times; (iii) Monte Carlo simulation studies of covariate-adjustment efficiency, sequential-design operating characteristics, and confidence-interval coverage under non-proportional hazards; (iv) cross-validated survival prediction with five machine-learning models; (v) a Bayesian Weibull analysis with skeptical, weakly-informative, and enthusiastic priors and sequential monitoring; and (vi) doubly robust estimation of heterogeneous absolute treatment effects. Levamisole plus 5-FU reduced the DFS hazard by 38% (adjusted HR = 0.621, 95% CI 0.496–0.776), gained 16.8 percentage points of five-year DFS and 0.63 disease-free years in restricted mean survival time. Under the redesigned monitoring plan the trial would have stopped for efficacy at the first interim analysis (118 events; HR 0.52, $Z = 3.38$ against a boundary of 2.96), and Bayesian monitoring would have signalled benefit at 25% of the information. Covariate adjustment raised simulated power from 72.6% to 80.9% for a moderate effect and directly illustrated the non-collapsibility of the hazard ratio, while the collapsible RMST difference was unchanged by adjustment. All code, data, and outputs are reproducible end-to-end from the accompanying repository.

Keywords: randomized clinical trial; survival analysis; group-sequential design; multiple imputation; competing risks; restricted mean survival time; Bayesian inference; machine learning; heterogeneous treatment effects.

1 Introduction

Colorectal cancer is one of the most common malignancies worldwide, and for patients undergoing potentially curative resection the question of whether systemic adjuvant chemotherapy should follow surgery remained open for much of the twentieth century. The North Central Cancer Treatment Group (NCCTG) program addressed this question directly: building on an earlier randomized comparison of observation against levamisole alone (Laurie et al., 1989), the investigators added a third arm combining the immunomodulator levamisole with the antimetabolite 5-fluorouracil (5-FU) (Moertel et al., 1990). The combination reduced recurrence and mortality substantially and became a standard of care for resected stage B/C colon

carcinoma. The individual-patient data from this program were subsequently contributed to the `survival` package for R (Therneau, 2026), where they now serve as one of the canonical datasets for teaching survival analysis; the version analysed here contains 929 patients from a single study.

Re-analyzing an old trial with new methods is more than an academic exercise, and five considerations motivated each component of this project. First, endpoint definitions and analysis conventions have evolved: modern adjuvant colorectal trials define disease-free survival (DFS) as the time to recurrence *or* death, whereas the published coding of this dataset treats death without recurrence as censoring; a rigorous re-analysis must therefore address both views and, importantly, the competing-risks consequences of the difference. Second, missing-data handling, covariate-adjusted estimation, and absolute-benefit reporting have all been refined since the original publications, and regulatory guidance now actively encourages pre-specified covariate adjustment in the primary analysis of randomized trials (U.S. Food and Drug Administration, 2023; Kahan et al., 2014; Sterne et al., 2009). Third, the trial-design landscape has changed: group-sequential designs with alpha- and beta-spending functions are now the default in oncology, and it is natural to ask whether this trial—run with a fixed-sample mindset—would have stopped early under a modern monitoring plan (Lan and DeMets, 1983; O’Brien and Fleming, 1979). Fourth, machine-learning prediction models and Bayesian machinery offer perspectives on prognosis and evidence accumulation that complement the classical toolkit (Harrell, 2015; Gelman et al., 2013). Finally, the modern estimand framework and recent methods for heterogeneous treatment effects invite a re-examination of *for whom* the benefit is largest (Hernán and Robins, 2020).

Contributions

This paper integrates all of these perspectives into a single coherent, fully scripted pipeline and makes the following contributions:

1. **A comprehensive treatment-effect re-estimate** on the composite DFS endpoint, combining multiple imputation, proportional-hazards diagnostics, competing-risk (Aalen–Johansen and Fine–Gray) sensitivity analyses, and absolute-benefit measures—five-year risk differences, numbers needed to treat, and restricted mean survival times (RMST), the latter providing a collapsible, assumption-light summary that directly contrasts with the non-collapsibility of the hazard ratio (Royston and Parmar, 2011; Uno et al., 2014).
2. **A counterfactual trial redesign and monitoring re-enactment**: the trial is re-sized and re-designed as a three-look group-sequential study (O’Brien and Fleming, 1979; Lan and DeMets, 1983; Hwang et al., 1990), and the monitoring plan is replayed on the real event stream, showing that a modern design would have stopped the trial for efficacy at the first interim analysis, with conditional-power and Bayesian monitoring cross-checks.
3. **Simulation-calibrated operating characteristics** for the design and analysis choices: the efficiency gain from covariate adjustment (including a direct quantitative demonstration of hazard-ratio non-collapsibility), the type-I error and stopping probabilities of the sequential design, and confidence-interval coverage under non-proportional hazards.
4. **An honest validation study of five survival prediction models** spanning penalized regression, kernel ensembles, and gradient boosting, quantifying the optimism gap between apparent and cross-validated discrimination.
5. **A Bayesian re-analysis** with explicit prior sensitivity (skeptical, weakly-informative, enthusiastic), event-driven sequential monitoring, posterior predictive checks, and baseline-model criticism by approximate leave-one-out cross-validation (Vehtari et al., 2017).

6. **A heterogeneous-treatment-effect analysis** on the absolute five-year risk scale, combining jackknife pseudo-observation regression (Andersen et al., 2003; Parner and Andersen, 2010), cross-fitted G-computation, and doubly robust AIPW estimation (Hernán and Robins, 2020), which locates the largest absolute benefits in the highest-risk patients.

Beyond the specific findings, the paper is deliberately methods-integrative: every paradigm is applied to the same data, under a common adjustment set and common software-engineering discipline (fixed seeds, file-based module interfaces, one-command reproduction), so that the contrasts between paradigms—collapsible versus non-collapsible estimands, alpha-spending versus posterior monitoring, apparent versus honest prediction accuracy—are attributable to the methods rather than to implementation differences.

The remainder of the paper is organized as follows. Section 2 describes the data and endpoint construction; Section 3 details all statistical methods; Section 4 presents results in five thematic blocks; and Section 5 discusses strengths, limitations, and conclusions.

2 Data and Endpoints

2.1 Source and patients

The analysis uses the individual-patient data of the NCCTG adjuvant colon cancer trial as distributed with the R package **survival** (Therneau, 2026): 929 patients with resected stage B/C colon carcinoma, all from a single study, randomized 1:1:1 to Observation ($n = 315$), Levamisole ($n = 310$), or Levamisole plus 5-FU ($n = 304$). The recorded baseline covariates are age, sex, bowel obstruction, perforation, adherence to nearby organs, number of positive lymph nodes, tumour differentiation, extent of local spread, and the surgery-to-registration interval. Randomization produced well-balanced arms: across eleven baseline variables the only nominal imbalance was sex (53.6% female in Lev+5FU vs. 47.3% in Observation; three-arm global exact test $p = 0.027$, pairwise $p = 0.13$), unremarkable among eleven simultaneous tests and not surviving multiplicity correction; pairwise standardized mean differences never exceeded $|0.216|$, and the overall picture is consistent with adequate randomization concealment.

2.2 Endpoint construction

The raw dataset stores two records per patient: **etype** = 1 records time to tumour recurrence with death treated as censoring, and **etype** = 2 records overall survival. Because current adjuvant colorectal conventions define DFS as recurrence *or* death from any cause, the data-curation module constructs the composite DFS endpoint explicitly (506 events; 468 recurrences plus 38 deaths without recorded recurrence), while overall survival (452 deaths) serves as the secondary endpoint and the recurrence-only coding with death as a competing event (38 competing events) is retained for sensitivity analyses. Covariate values are verified to be identical across the two records of each patient; event times satisfy the required orderings (every recurrence time does not exceed the corresponding death time, and deaths without recurrence share the death time). Median follow-up, estimated by reverse Kaplan–Meier, is 6.4 years (maximum 9.1 years).

2.3 Missing data

Two covariates have missing values: number of positive lymph nodes (18 patients, 1.9%) and tumour differentiation (23 patients, 2.5%); 95.6% of patients are complete cases. Missing values are handled by multiple imputation as described below, with complete-case analyses reported for comparison.

Table 1: Baseline characteristics by randomized arm (percentages of the full arm; nodes and differentiation include a small missing category).

Characteristic	Observation	Levamisole	Levamisole + 5-FU
Randomized (n)	315	310	304
Age, years, mean (SD)	59.5 (12.0)	60.1 (11.6)	59.7 (12.3)
Male sex, n (%)	166 (52.7)	177 (57.1)	141 (46.4)
Obstruction, n (%)	63 (20.0)	63 (20.3)	54 (17.8)
Perforation, n (%)	9 (2.9)	10 (3.2)	8 (2.6)
Adherence to organs, n (%)	47 (14.9)	49 (15.8)	39 (12.8)
Positive nodes, median [IQR]	2 [1–5]	2 [1–5]	2 [1–4]
≥ 4 positive nodes, n (%)	119 (37.8)	105 (33.9)	92 (30.3)
Poor differentiation, n (%)	52 (16.5)	44 (14.2)	54 (17.8)
Serosal/contiguous extent, n (%)	269 (85.4)	271 (87.4)	262 (86.2)
Long surgery-to-registration, n (%)	224 (71.1)	230 (74.2)	228 (75.0)
Composite DFS events, n (%)	190 (60.3)	182 (58.7)	134 (44.1)
Deaths, n (%)	168 (53.3)	161 (51.9)	123 (40.5)

3 Methods

3.1 Frequentist efficacy analysis

Disease-free survival and overall survival were summarized by Kaplan–Meier curves with point-wise 95% confidence intervals (log–log transformation) and compared with log-rank tests, both overall (2 degrees of freedom) and pairwise between arms. Treatment effects were estimated with Cox proportional-hazards models fitted by Efron’s method for ties. Three specifications are reported for the primary endpoint: an unadjusted model; a covariate-adjusted complete-case model; and the pre-specified covariate-adjusted model fitted on each of the $m = 20$ multiply imputed datasets and pooled by Rubin’s rules (Rubin, 1987). The imputation models used chained equations with the Nelson–Aalen cumulative hazard and the event indicator as predictors, following White and Royston (2009) and van Buuren and Groothuis-Oudshoorn (2011), so that imputation is compatible with the subsequent Cox analysis. The adjustment set comprised treatment arm, age, sex, obstruction, perforation, adherence, number of positive nodes, differentiation, extent of local spread, and surgery-to-registration interval; an analogous pooled model was fitted for overall survival.

The proportional-hazards assumption was examined with scaled Schoenfeld residuals (Grambsch and Therneau, 1994), both graphically (loess smooths against time) and with the global chi-square test. Two pre-specified sensitivity analyses address potential violations and endpoint-definition concerns: first, flexible functional forms for age and nodes via restricted penalized splines (4 degrees of freedom each), compared with the linear model by likelihood-ratio test; second, competing-risk analyses on the recurrence endpoint, contrasting the Aalen–Johansen cumulative incidence function (Aalen, 1978) with the naive one-minus-Kaplan–Meier estimate, and fitting Fine–Gray subdistribution-hazard models (Fine and Gray, 1999) adjusted for the same covariate set.

3.2 Absolute benefit and restricted mean survival time

Because hazard ratios are not directly interpretable as absolute benefit, five-year Kaplan–Meier survival probabilities with 95% confidence intervals were computed by arm for both endpoints, along with between-arm risk differences and numbers needed to treat (NNT) whose 95% confidence intervals were obtained by nonparametric bootstrap with 2,000 resamples.

In addition, the primary comparison was re-expressed through the restricted mean survival time (Royston and Parmar, 2011; Uno et al., 2014): the area under the Kaplan–Meier survival curve up to a horizon τ , estimated at $\tau = 3, \dots, 8$ years with percentile bootstrap confidence intervals (2,000 resamples), and summarized both as a between-arm difference (disease-free years gained) and as a ratio. Unlike the hazard ratio, the RMST difference is collapsible and requires no proportional-hazards assumption. A covariate-adjusted RMST difference at $\tau = 5$ years was obtained by regressing leave-one-out jackknife pseudo-observations of RMST(5) (Andersen et al., 2003; Parner and Andersen, 2010) on treatment and the standardized adjustment set, with bootstrap standard errors; because the estimand is collapsible, the adjusted and unadjusted analyses target the same quantity, providing a direct empirical contrast with the non-collapsibility of the hazard ratio quantified in the simulation study below. RMST differences were also computed within the obstruction stratum, the covariate with the strongest proportional-hazards violation.

3.3 Subgroup and interaction analyses

Seven clinically motivated subgroups were pre-specified for the primary comparison of levamisole plus 5-FU against observation: age (< 60 vs. ≥ 60 years), sex, nodal burden (< 4 vs. ≥ 4 positive nodes), obstruction, local extent (submucosa/muscle vs. serosa/contiguous), differentiation (well/moderate vs. poor), and surgery-to-registration interval. Within each subgroup level, unadjusted Cox models estimated subgroup-specific hazard ratios; effect modification was tested by Wald tests of treatment-by-subgroup interaction terms in models adjusted for the nodal count. Given seven interactions tested, only p -values below a Bonferroni-adjusted threshold of $0.05/7 \approx 0.0071$ are interpreted as evidence of modification.

3.4 Trial redesign and re-enacted monitoring

To redesign the trial under modern standards, planning assumptions were taken from the re-analysis itself. The control arm’s five-year DFS of 42.4% and median DFS of 2.96 years imply an exponential baseline hazard $\lambda_0 = \log 2/2.96 = 0.234$ per year; the clinically meaningful target was set at a hazard ratio of 0.65, close to the effect actually observed. A fixed one-sided superiority design at $\alpha = 0.025$ with 90% power, 1:1 allocation, three years of uniform accrual and a minimum of five years of additional follow-up, requires 226 DFS events and approximately 322 patients (Schoenfeld’s formula (Schoenfeld, 1983), verified with `gsDesign::nSurv` (Anderson, 2026)).

The modern redesign is a three-look group-sequential trial with analyses after 50% and 75% of the planned events and at completion. Efficacy boundaries use O’Brien–Fleming alpha-spending (O’Brien and Fleming, 1979; Lan and DeMets, 1983); futility boundaries use a non-binding Hwang–Shih–DeCani beta-spending with shape parameter -4 (Hwang et al., 1990), which yields aggressive early stopping for futility under the null while allowing the sponsor to disregard the futility boundary without inflating the type-I error. The resulting maximum sample size is 235 events (approximately 336 patients), a 4.1% inflation over the fixed design.

The monitoring plan was then *re-enacted on the real trial data*: because the public dataset contains no calendar dates, an information-time snapshot of the interim stage was reconstructed by administratively censoring every patient’s follow-up at the time (years since randomization) at which the 118th DFS event of the actual trial was observed, refitting the Cox model on the censored-at-interim snapshot, and comparing the resulting Z -statistic with the design boundaries. Conditional power for the final analysis was computed under three assumptions for the remaining drift (null effect, design alternative, and the observed interim trend), with the sign convention that positive Z favours the experimental arm; the hand-rolled calculation was cross-checked with `gsDesign`’s conditional-power machinery.

3.5 Simulation studies

Three Monte Carlo studies quantify the operating characteristics behind the design and analysis choices. The data-generating mechanism was calibrated to the trial itself: covariates were drawn by bootstrap resampling of the complete-case patients, the prognostic index used hazard-ratio coefficients from the adjusted Cox model, event times were exponential with hazard $\lambda_0 \exp(\text{lp}_i + \log(\text{HR}) \cdot \text{treat}_i)$, accrual times were uniform on three years, and administrative censoring occurred at eight years.

Study A (covariate-adjustment efficiency) simulated 1,000 trials per scenario for a two-arm comparison with $n = 322$ at hazard ratios of 1.0 (type-I error) and 0.72 (power), crossing three levels of prognostic strength obtained by scaling the calibrated covariate coefficients by 0, 0.5, and 1 (standard deviations of the prognostic index: 0, 0.208, and 0.416 on the log-hazard scale). Each replicate fitted both the unadjusted and the fully adjusted Cox model, recording estimates, model-based standard errors, empirical standard deviations, and rejection frequencies. The contrast between the unadjusted estimator and the data-generating conditional hazard ratio quantifies the non-collapsibility of the hazard ratio (Martinussen and Vansteelandt, 2013).

Study B (sequential-design operating characteristics) simulated the full monitoring procedure of Section 3.4 (3,000 trials per scenario) with event-triggered looks and calendar-time censoring, under the alternative ($\text{HR} = 0.65$) and under the null with the futility boundary adhered to and ignored.

Study C (interval estimation under non-proportional hazards) simulated 500 trials with a piecewise-constant log-hazard ratio (HR 0.45 in year one, HR 0.80 thereafter; the exponential memorylessness property makes the piecewise construction exact), comparing the coverage and width of Wald and percentile-bootstrap confidence intervals for the log hazard ratio (100 bootstrap resamples per replicate) against the simulation estimand.

3.6 Machine-learning prediction

Five prognostic models for the composite DFS endpoint were compared with repeated (5×5 -fold) stratified cross-validation: a clinical Cox model with main effects (benchmark); an elastic-net penalized Cox model; a random survival forest; gradient boosting with Cox partial-likelihood loss; and XGBoost with the Cox objective. Treatment assignment was deliberately excluded from all predictors, so the models quantify baseline risk only—a task distinct from treatment-effect estimation. Discrimination was assessed with Harrell’s C-index and the time-dependent IPCW AUC at five years, prediction error with the integrated Brier score over years 1–5 (Graf et al., 1999), and calibration by deciles of predicted risk against Kaplan–Meier observed risk. The apparent (in-sample) performance of each model was also computed to quantify the optimism gap. Permutation importance (20 shuffles per covariate) interpreted the winning model.

3.7 Bayesian re-analysis

A Weibull proportional-hazards model was fitted by NUTS Hamiltonian Monte Carlo (PyMC; Salvatier et al., 2016) to the primary two-arm comparison, with right censoring handled by likelihood factorization (the censored observations contribute their survival function). The parametrization $H(t | x) = (t/\beta_0)^\alpha \exp(\eta_i)$, with η_i the linear predictor, makes $\exp(\beta_{\text{treat}})$ the hazard ratio. Three priors for the treatment log-hazard ratio were compared: *skeptical* ($\mathcal{N}(0, 0.421^2)$, placing 5% prior mass below $\text{HR} = 0.5$); *weakly-informative* ($\mathcal{N}(0, 1)$); and *enthusiastic* ($\mathcal{N}(\log 0.65, 0.421^2)$, centred at the design alternative). Sequential monitoring replayed evidence accumulation at 25%, 50%, 75%, and 100% of the observed DFS events under the skeptical prior; posterior predictive checks replicated the control arm’s Kaplan–Meier curve

from 100 posterior draws; and approximate leave-one-out cross-validation (Vehtari et al., 2017) compared the Weibull baseline against the exponential special case. Convergence was assessed by $\hat{R}$ and effective sample sizes.

3.8 Heterogeneous treatment effects

Heterogeneity was characterized on the absolute five-year DFS event risk. Jackknife pseudo-observations for the five-year risk (Andersen et al., 2003; Parner and Andersen, 2010) converted the censored outcome into a regression-ready continuous outcome. The population-average effect was estimated three ways: the unadjusted Kaplan–Meier difference with bootstrap confidence intervals; cross-fitted G-computation; and a cross-fitted doubly robust AIPW estimator with influence-function standard errors (Hernán and Robins, 2020), using the empirical randomization probability as the propensity. Conditional (patient-level) effects were estimated by a cross-fitted T-learner (penalized outcome models by arm), summarized by the distribution and quartiles of predicted absolute benefit and by the dependence of benefit on baseline risk.

3.9 Software and reproducibility

The pipeline consists of six R modules (data curation, descriptives, frequentist analysis, subgroups, trial design, simulation) and four Python modules (machine learning, Bayesian inference, heterogeneity, RMST), orchestrated by a single `Makefile` invocation. All scripts use fixed seeds (global seed 20260916), communicate exclusively through files in `data/` and `outputs/`, and embed quality-control assertions (record counts, within-patient covariate consistency, end-point logic). Figures are produced as vector PDF for the paper and 300-dpi PNG for the repository README. Exact package versions are archived with the repository, and every number in this paper was regenerated and verified from the committed outputs during preparation of the manuscript.

4 Results

4.1 Treatment effects on disease-free and overall survival

Kaplan–Meier analysis

Figures 1 and 2 display Kaplan–Meier estimates for the composite DFS endpoint and for overall survival, each with the numbers at risk. The three curves separate early and the separation widens over the first three years, driven almost entirely by the levamisole-plus-5-FU arm. The overall log-rank test is highly significant for DFS ($\chi^2 = 21.0$ on 2 degrees of freedom, $p < 0.001$); pairwise comparisons show that levamisole alone is indistinguishable from observation (HR 0.968, $p = 0.76$), whereas levamisole plus 5-FU is superior to both observation (HR 0.623, $p < 0.001$) and to levamisole alone ($p < 0.001$). The pattern for overall survival is similar but attenuated (adjusted HR 0.688, 95% CI 0.544–0.871), consistent with post-recurrence salvage therapies diluting survival differences.

Cox models

Table 2 assembles the treatment-effect estimates across specifications. Adjustment for the ten baseline covariates, pooled over 20 multiple imputations by Rubin’s rules, moves the primary contrast only marginally (unadjusted HR 0.623 vs. adjusted 0.621), as expected under randomization with balanced arms; the complete-case analysis agrees closely. The adjusted covariate effects themselves (Figure 3) are clinically coherent: each additional positive lymph node multiplies the hazard by 1.08 (95% CI 1.06–1.10), spread to contiguous structures triples




Figure 1: Kaplan–Meier estimates of the composite disease-free survival endpoint (primary) with pointwise 95% confidence bands and numbers at risk.




Figure 2: Kaplan–Meier estimates of overall survival (secondary endpoint) with pointwise 95% confidence bands and numbers at risk.

the hazard relative to submucosal involvement (HR 3.19, 95% CI 1.32–7.75), obstruction raises it by 25%, and a long surgery-to-registration interval is protective in this cohort (HR 0.76, 95% CI 0.62–0.92).

Table 2: Treatment-effect estimates for the composite DFS endpoint across analysis specifications. The Fine-Gray model targets the recurrence endpoint with death without recurrence as a competing event.

Analysis	Contrast	HR	95% CI	<i>p</i>
Unadjusted Cox	Lev vs. Obs	0.968	0.790–1.187	0.757
Unadjusted Cox	Lev+5FU vs. Obs	0.623	0.499–0.777	< 0.001
Adjusted Cox, MI ($m = 20$)	Lev vs. Obs	0.947	0.770–1.165	0.607
Adjusted Cox, MI ($m = 20$)	Lev+5FU vs. Obs	0.621	0.496–0.776	< 0.001
Adjusted Cox, MI (OS endpoint)	Lev+5FU vs. Obs	0.688	0.544–0.871	0.002
Fine-Gray subdistribution	Lev+5FU vs. Obs	0.596	0.469–0.756	< 0.001




Figure 3: Adjusted hazard ratios for disease-free survival from the multiply imputed Cox model (log scale). Red indicates estimates whose 95% confidence interval excludes 1.

Absolute benefit

Five-year Kaplan–Meier DFS probabilities were 42.4% (95% CI 36.9–47.8) under observation, 44.2% (38.6–49.6) under levamisole, and 59.2% (53.4–64.5) under levamisole plus 5-FU; the corresponding overall-survival probabilities were 52.6%, 53.5%, and 63.4%. The bootstrap 95% CI for the five-year DFS risk difference of levamisole plus 5-FU against observation is –16.7 percentage points (–24.2 to –9.4), corresponding to a number needed to treat of six patients to

prevent one additional DFS event at five years. Levamisole alone shows no detectable absolute benefit (-1.8 points, -9.3 to $+5.7$).

Proportional-hazards diagnostics

The scaled Schoenfeld-residual analysis flags two covariates: tumour differentiation ($p < 0.001$) and obstruction ($p = 0.008$), with a significant global test ($\chi^2 = 42.2$, 14 df, $p < 0.001$). Reassuringly, the treatment indicator itself shows no evidence of a time-varying effect ($p = 0.51$), so the primary hazard-ratio summary is defensible. Figure 4 shows the residual patterns; the obstruction effect visibly weakens after year two. Replacing linear age and nodes with 4-degree-of-freedom penalized splines significantly improves the log-likelihood (likelihood-ratio $p < 0.001$) without affecting the treatment estimate.




Figure 4: Proportional-hazards diagnostics: scaled Schoenfeld residuals with loess smooths for treatment, positive nodes, and obstruction (which exhibits a significant violation, $p = 0.008$).

Competing-risk sensitivity analysis

Thirty-eight patients died without a recorded recurrence, and treating these deaths as censoring—as the published coding does—formally assumes they carry no recurrence-risk information. The competing-risk view shows the practical consequence (Figure 5): the Aalen–Johansen cumulative incidence of recurrence sits below the naive one-minus-Kaplan–Meier curve for every arm, with the gap reaching about one percentage point at eight years in the observation arm. The Fine–Gray subdistribution model estimates the levamisole-plus-5-FU effect on the recurrence subdistribution hazard at 0.596 (95% CI 0.469–0.756)—virtually

identical to the cause-specific reading of 0.623—so the conclusion of a substantial treatment benefit is robust to the endpoint-definition choice.

Cumulative incidence of recurrence: competing-risk vs naive analysis




Figure 5: Cumulative incidence of recurrence: Aalen-Johansen estimates treating death without recurrence as a competing event (solid) versus the naive one-minus-Kaplan-Meier estimator that censors such deaths (dashed).

Restricted mean survival time

The RMST analysis restates the benefit in time units without any proportional-hazards assumption (Table 3, Figure 6). Within a five-year horizon, patients randomized to levamisole plus 5-FU can expect 3.57 disease-free years on average versus 2.94 under observation—a gain of 0.63 years (95% CI 0.33 to 0.92), roughly seven and a half months, with an RMST ratio of 1.21 (1.11 to 1.33). The gain widens with the horizon, reaching 1.13 years (0.62 to 1.65) by eight years as the curves continue to separate. The covariate-adjusted pseudo-observation regression yields 0.57 years (0.28 to 0.85), closely agreeing with the unadjusted estimate—as it must, because the RMST difference, unlike the hazard ratio, is collapsible: covariate adjustment changes precision, not the estimand. Within the obstruction stratum—the covariate with the strongest proportional-hazards violation ($p = 0.008$)—the five-year RMST difference remains substantial (0.60 years, 0.27 to 0.94, without obstruction; 0.83 years, 0.11 to 1.55, with obstruction), confirming that the treatment conclusion does not rest on the proportional-hazards assumption.

4.2 Subgroups, trial redesign, and re-enacted monitoring

Subgroup analyses

Figure 7 presents the seven pre-specified subgroup analyses for the primary contrast. The hazard ratio is remarkably stable across subgroups—every point estimate lies between 0.46 and 0.81 and every confidence interval overlaps the overall estimate of 0.62. Only one nominal interaction reaches $p < 0.05$: sex (interaction $p = 0.043$), with a substantially larger effect

Table 3: Restricted mean survival time for the primary comparison (Lev+5FU vs. observation, DFS endpoint). Difference = disease-free years gained; percentile bootstrap CIs (2,000 resamples).

Horizon τ	RMST Obs (y)	RMST Lev+5FU (y)	Difference (95% CI)	Ratio (95% CI)
3	2.03	2.34	0.31 (0.14–0.47)	1.15 (1.07–1.24)
4	2.50	2.96	0.46 (0.23–0.68)	1.18 (1.09–1.29)
5	2.94	3.57	0.63 (0.33–0.92)	1.21 (1.11–1.33)
6	3.36	4.14	0.79 (0.42–1.15)	1.23 (1.12–1.37)
7	3.75	4.70	0.95 (0.51–1.39)	1.25 (1.13–1.40)
8	4.12	5.24	1.13 (0.62–1.65)	1.27 (1.14–1.43)




Figure 6: Restricted mean survival time analysis. Left: the RMST is the shaded area under the Kaplan–Meier curve up to the horizon $\tau = 5$ years (shown for the observation arm). Right: RMST difference (disease-free years gained with levamisole plus 5-FU) with 95% bootstrap confidence intervals as a function of the horizon.

in men (HR 0.46, 95% CI 0.33–0.64) than in women (HR 0.79, 95% CI 0.58–1.07). Against the Bonferroni-adjusted threshold of 0.0071 appropriate for seven tests, this signal does not survive, and no biological mechanism has been established for sex-specific efficacy of 5-FU-based regimens in this setting. The honest reading is *consistent relative effect across all examined subgroups*.

Subgroup analyses: disease-free survival, Lev+5-FU vs observation

Unadjusted Cox models within subgroups; interaction tests adjusted for nodal count.
Right: hazard ratios with 95% CIs and interaction p-values.




Figure 7: Pre-specified subgroup analyses for disease-free survival, levamisole plus 5-FU versus observation. Right column: hazard ratios with 95% CIs and interaction-test *p*-values (adjusted for nodal count).

The redesigned trial

Table 4 summarizes the fixed design and the group-sequential redesign. For the target hazard ratio of 0.65 at one-sided $\alpha = 0.025$ with 90% power, the fixed design requires 226 DFS events (approximately 322 patients under the assumed accrual and follow-up). The three-look sequential design inflates the maximum to 235 events but compensates with expected savings whenever the effect is real or absent. The O'Brien–Fleming efficacy boundary is naturally conservative early ($Z = 2.963$ at the first interim) and relaxes toward 2.014 at the final analysis; the non-binding futility boundary starts at $Z = 0.079$, meaning a first look with essentially no observed trend would trigger the futility signal. Power falls steeply for weaker effects: at $HR = 0.80$, nearly 850 events would be required—a useful reminder of why ambitious trials with modest expected effects grow so large.

Table 4: The redesigned group-sequential trial: boundaries and operating characteristics from `gsDesign`. Efficacy uses O’Brien–Fleming alpha-spending; futility uses a non-binding Hwang–Shih–DeCani beta-spending ($\gamma = -4$).

Analysis	Events	Effic. Z	1-sided p	HR at bound	Futil. Z	HR at bound
Interim 1 (50%)	118	2.963	0.0015	0.579	0.078	0.986
Interim 2 (75%)	177	2.359	0.0091	0.701	1.006	0.859
Final (100%)	235	2.014	0.0220	0.769	2.014	0.769

Design-theoretic probabilities: $P(\text{reject} \mid \text{HR} = 0.65) = 0.900$; $P(\text{reject} \mid \text{HR} = 1) = 0.024$ (futility adhered)




Figure 8: Group-sequential boundaries of the redesigned trial on the Z scale (left axis) and one-sided p -value scale (right axis). Red: efficacy (O’Brien–Fleming spending); blue: non-binding futility (Hwang–Shih–DeCani).

Re-enacting interim monitoring on the real data

The most striking result of this re-analysis emerges when the redesigned monitoring plan is applied to the *actual* event stream. Censoring every patient’s follow-up at the time at which the 118th DFS event was observed (information fraction 50%) and refitting the two-arm Cox model yields an interim hazard ratio of 0.52 and a Z -statistic of 3.38—*above the efficacy boundary of 2.963* (Figure 9). Under the redesign, the trial would have stopped at the first interim analysis and declared the levamisole plus 5-FU regimen superior to observation, after only 118 of the 506 events that were eventually observed. Conditional power calculations quantify how decisive this interim result is: the probability of eventually crossing the final boundary is 70.6% even under the null, 99.8% under the design alternative, and essentially 100% if the observed interim trend persists. A cross-check with `gsDesign`’s conditional-power machinery returns somewhat higher values (0.834 and 0.999 under the null and alternative), because it accumulates the crossing probabilities at *both* remaining analyses with boundaries re-optimized after the interim, whereas the hand-rolled calculation uses a single-look normal approximation to the final analysis only. The contrast with the Bayesian view is instructive: sequential posterior probabilities under the skeptical prior (Section 4.4) would have signalled benefit even earlier, at the 25% information look.

Two interpretive cautions apply. First, the historical trial was not designed with this monitoring plan, so the re-enactment is a counterfactual exercise, not a critique of the original conduct; its message is about how much evidence was already present early in the event stream. Second, the interim hazard ratio of 0.52 is more extreme than the final estimate of 0.62, a reminder that early random highs in Z are partly regression artefacts; the design’s conservative early boundary exists precisely to price this randomness in.




Figure 9: Re-enacted interim monitoring using the real trial data. The observed interim result after 118 events (HR 0.52, $Z = 3.38$, green marker) crosses the O’Brien–Fleming efficacy boundary (2.96), so the redesigned trial would have stopped for efficacy at the first interim analysis.

4.3 Simulation studies

Study A: covariate adjustment buys power

Table 5 and Figure 10 report 1,000 simulated trials per scenario at $n = 322$ and target hazard ratio 0.72. With no prognostic covariates, adjustment neither helps nor hurts, as theory predicts: power is 71.6% unadjusted versus 70.4% adjusted, and the type-I error is 2.5% and 2.9% respectively (2.2% and 3.2% at full strength)—all within Monte Carlo error of the nominal 2.5%. As prognostic strength grows to the trial-calibrated value (SD of the prognostic index 0.416 on the log-hazard scale), adjusted power rises to 80.9% against 72.6% unadjusted—an efficiency gain equivalent to roughly 70 additional events (about a 23% larger event count), obtained at zero cost in bias or type-I error. This is the quantitative case behind current regulatory encouragement of pre-specified covariate adjustment (U.S. Food and Drug Administration, 2023; Kahan et al., 2014).

The bias column also teaches a subtler lesson: at full prognostic strength, the unadjusted estimator averages +0.033 on the log scale relative to the data-generating conditional hazard ratio—that is, an unadjusted marginal HR of about 0.745 when the conditional truth is 0.72. This attenuation toward the null is the hazard ratio’s non-collapsibility (Martinussen and Vansteelandt, 2013): unlike a linear-model coefficient or a risk difference (or, as demonstrated above, the RMST difference), the marginal HR changes when strong prognostic covariates are removed from the model even without any confounding. It explains why adjusted and unadjusted estimates in randomized trials need not coincide, and why the near-coincidence of 0.623 and 0.621 in our real data reflects the modest prognostic strength of the adjustment set rather than a mathematical necessity.

Table 5: Simulation Study A: covariate-adjustment efficiency (1,000 replicates per scenario, $n = 322$, one-sided $\alpha = 0.025$). Power at HR = 0.72; type-I error at HR = 1. “Bias” contrasts each estimator with the conditional log-HR used to generate the data.

Prognostic strength	HR	Power (unadj.)	Power (adj.)	Bias (unadj.)	Bias (adj.)
None (SD 0.000)	1.00	0.025	0.029	−0.006	−0.006
Moderate (SD 0.208)	1.00	0.025	0.025	+0.004	+0.005
Full (SD 0.416)	1.00	0.022	0.032	+0.000	−0.001
None (SD 0.000)	0.72	0.716	0.704	+0.000	−0.009
Moderate (SD 0.208)	0.72	0.725	0.766	+0.012	−0.006
Full (SD 0.416)	0.72	0.726	0.809	+0.033	−0.010

Study B: operating characteristics of the sequential design

Simulating the full monitoring procedure (3,000 trials per scenario, event-triggered looks with calendar-time censoring) validates the design and quantifies the value of the futility boundary (Table 6, Figure 11). Under the alternative (HR 0.65), the design rejects in 82.9% of trials against the design-theoretic 90%. Three mechanisms account for the shortfall, all instructive. First, the simulated primary analysis is the *unadjusted* Cox test, whose target is the marginal hazard ratio—attenuated toward the null by the strong calibration covariates through exactly the non-collapsibility quantified in Study A. Second, futility adherence stops 6.8% of trials under the alternative that might eventually have crossed. Third, event-triggered looks with discrete calendar-time cuts approximate, but do not reproduce, the continuous information accrual assumed by the spending functions. Stopping is dominated by the second interim (42.2%) rather than the first (18.6%), and the expected event count is 183.4 versus 235 at completion: the design saves, on average, 51.6 events (roughly 22%) when the drug works.




Figure 10: Simulation Study A: empirical power of the unadjusted versus covariate-adjusted Cox test as a function of prognostic strength (target HR 0.72, $n = 322$, 1,000 replicates per scenario).

Under the null with futility adhered to, the type-I error is 2.4% and the trial stops for futility at the first interim in 54.9% of cases. Ignoring futility (the non-binding option) inflates the type-I error only trivially, to 2.7%, confirming the design’s stated robustness.

Table 6: Simulation Study B: simulated operating characteristics of the group-sequential design (3,000 trials per scenario).

Scenario	Eff. IA1	Eff. IA2	Final	Total eff.	Exp. events	Events saved
HR = 0.65 (as designed)	0.186	0.422	0.222	0.829	183.4	51.6
HR = 1, futility adhered	0.002	0.008	0.014	0.024	152.7	82.4
HR = 1, futility ignored	0.001	0.009	0.017	0.027	234.4	0.6

Study C: interval estimation under non-proportional hazards

When the true effect is time-varying (HR 0.45 in year one, weakening to 0.80 thereafter) and the analysis fits a single proportional-hazards Cox model, the estimator targets a weighted average of the two regimes (simulation estimand: $\log\text{-HR} = -0.407$, i.e. a pooled HR of about 0.666). Over 500 replicates, the Wald interval achieved 92.6% coverage of this estimand and the percentile bootstrap 90.4%—both marginally below their nominal 95% (Table 7). The bootstrap is slightly narrower (mean width 0.434 vs. 0.443) but not more accurate. The lesson is double-edged: neither standard interval collapses under moderate non-proportional hazards—the practical damage is a shifted estimand rather than drastically wrong uncertainty—yet the modest undercoverage confirms that misspecification costs are real, reinforcing the value of the diagnostics and of the RMST analysis of Section 4.1.

Simulation Study B: simulated stopping probabilities of the group-sequential design

3000 simulated trials per scenario; bars show probability of stopping at each analysis by reason



Figure 11: Simulation Study B: simulated stopping probabilities by analysis and reason under the alternative (right) and under the null with futility adherence (left).

Table 7: Simulation Study C: 95% interval performance under a time-varying hazard ratio (500 replicates, $n = 335$; 100 bootstrap resamples per replicate).

Method	Coverage	Mean width	Mean log-HR
Cox Wald CI	0.926	0.443	-0.407
Bootstrap percentile	0.904	0.434	-0.407

Simulation Study C: interval estimates under non-proportional hazards

True effect is time-varying (HR 0.45 in year 1, HR 0.80 thereafter); dashed line = simulation estimand (500 replicate)




Figure 12: Simulation Study C: Wald versus bootstrap percentile intervals for the log hazard ratio across 500 simulated trials under a time-varying true effect. The dashed line marks the simulation estimand.

4.4 Machine-learning prediction and Bayesian inference

Cross-validated prediction performance

Table 8 and Figure 13 compare the five prognostic models over 25 cross-validation folds (888 complete cases). The random survival forest achieves the best discrimination (cross-validated $C = 0.646$, IPCW AUC at five years 0.677) and the lowest integrated Brier score (0.217 over years 1–5), ahead of the clinical Cox model ($C = 0.634$), gradient boosting (0.633), XGBoost (0.625), and the elastic-net Cox model (0.620). The spread among models is modest—about 0.026 in the C -index—and the honest uncertainty (fold-to-fold SD ≈ 0.025) cautions against declaring a definitive winner; what the comparison really shows is that, with ten baseline covariates and ~ 300 events, flexible learners extract only a little more signal than a well-specified clinical model.

The apparent-versus-cross-validated contrast in Table 8 quantifies overfitting: XGBoost’s apparent C of 0.741 drops to 0.625 under honest validation, a 0.12 optimism gap, versus a 0.07 gap for the random forest and essentially none for the clinical Cox model. This gap is the single most important column in the table—it is the difference between the performance one would report from a naive in-sample fit and the performance a new patient would actually experience. Calibration of the winning model (Figure 14) is satisfactory across the risk spectrum: within deciles of predicted five-year risk, the mean predicted risk tracks the Kaplan–Meier observed risk within 6.5 percentage points. Permutation importance is dominated by the positive lymph-node count (mean C -index drop 0.117 per permutation, an order of magnitude above every other covariate), followed at a distance by advanced local extent, female sex, and a long surgery-to-registration interval—a ranking that mirrors the adjusted Cox effects of Figure 3 and reinforces the clinical coherence of the learned model.

Table 8: Machine-learning prediction of the composite DFS endpoint (888 complete cases; 5×5 -fold stratified cross-validation). Treatment excluded from all predictors. AUC at five years is IPCW weighted; IBS is the integrated Brier score over years 1–5.

Model	CV C -index	CV AUC (5y)	CV IBS	Apparent C
Random survival forest	0.646	0.677	0.217	0.713
Clinical Cox (main effects)	0.634	0.652	0.222	0.651
Gradient boosting (Cox loss)	0.633	0.662	0.225	0.717
XGBoost (Cox objective)	0.625	–	–	0.741
Elastic-net Cox	0.620	0.646	0.226	0.628

Bayesian re-analysis: priors and posteriors

The Weibull model converged cleanly under all three prior scenarios (two chains, target acceptance 0.9; $\hat{R} \approx 1.00$ throughout). Table 9 and Figure 15 display the central result: the posterior hazard ratio is essentially insensitive to the prior, ranging only from 0.591 (weakly-informative) to 0.610 (skeptical). Under the skeptical prior—which places 95% of its mass between hazard ratios of about 0.44 and 2.28 and only 5% below 0.5—the posterior mean is 0.610 with 95% credible interval 0.485–0.757, $P(\text{HR} < 1) = 1.000$ to three decimals, and $P(\text{HR} < 0.8) = 0.994$. The data are simply strong enough to overwhelm any of the three priors, which is itself an informative finding: conclusions of substance do not require prior consensus.




Figure 13: Cross-validated discrimination by model. Left: distribution of the C -index across 25 folds. Right: IPCW AUC at five years with cross-validation standard errors.




Figure 14: Discrimination and calibration of the random survival forest. Left: ROC curve for five-year event classification. Right: predicted versus Kaplan–Meier observed five-year event risk by decile of predicted risk.

Table 9: Bayesian Weibull analysis of the primary contrast under three prior specifications (2 chains; 4,000 posterior draws each), with the frequentist adjusted estimate for comparison.

Prior	Mean (log-HR)	Post. HR	95% CrI	P(HR<1)	P(HR<0.8)
Skeptical $\mathcal{N}(0, 0.421^2)$	0	0.610	0.485–0.757	1.000	0.994
Weakly-inf. $\mathcal{N}(0, 1)$	0	0.591	0.466–0.738	1.000	0.997
Enthusiastic $\mathcal{N}(\log 0.65, 0.421^2)$	−0.431	0.592	0.474–0.734	1.000	0.997
Frequentist (adj. Cox, MI) [ref.]	–	0.621	0.496–0.776	–	–




Figure 15: Prior sensitivity of the Bayesian treatment-effect estimate. Left: the three prior densities on the hazard-ratio scale. Right: the corresponding posterior densities, with the frequentist point estimate (dotted) for reference.

Bayesian sequential monitoring

Figure 16 replays evidence accumulation under the skeptical prior at four event-driven looks. By the 25% look (79 events), the posterior mean hazard ratio is already 0.601 with $P(HR < 1) = 0.998$ and a credible interval excluding 1 (upper bound 0.88); from the 50% look (158 events) onward the probability is essentially 1. Under the common Bayesian monitoring rule of stopping when $P(HR < 1)$ exceeds 0.975, this trial would have stopped at the *first quarter* of the evidence—even earlier than the frequentist group-sequential re-enactment of Section 4.2, which required 50% of the information and a much larger Z . The contrast is partly definitional (the Bayesian posterior probability does not pay a multiplicity price for looks) and partly prior-driven (the skeptical prior pulls toward the null early on, making its rapid crossing more impressive); both mechanisms deserve emphasis when Bayesian designs are proposed as replacements for alpha-spending.




Figure 16: Bayesian sequential monitoring under the skeptical prior. Left: posterior mean hazard ratio with 95% credible intervals at each look. Right: the monitoring statistic $P(HR < 1)$ against the 0.975 stopping threshold.

Model criticism

The posterior predictive check (Figure 17) replicates the control arm’s Kaplan–Meier curve from 100 posterior draws: the observed curve lies comfortably within the 95% posterior predictive band throughout follow-up, with no systematic deviation in the tail where a poorly specified baseline hazard would typically betray itself. Approximate leave-one-out cross-validation strongly prefers the Weibull baseline over the exponential special case (ELPD difference -18.0 , all stacking weight on the Weibull), and the posterior shape parameter $\alpha = 0.765$ (95% CrI 0.695–0.837) confirms a genuinely decreasing baseline hazard—patients who survive the first two years have progressively lower event rates, exactly as clinical experience suggests.




Figure 17: Posterior predictive check for the control arm: replicated Kaplan–Meier curves from 100 posterior draws versus the observed curve.

4.5 Heterogeneous treatment effects

Average effects on the five-year absolute risk

Table 10 reports the average treatment effect of levamisole plus 5-FU versus observation on the five-year DFS event risk, estimated three ways on the jackknife pseudo-observations ($n = 594$ complete cases). The three estimators agree closely: the unadjusted Kaplan–Meier difference is -16.8 percentage points (bootstrap 95% CI -24.5 to -9.1), cross-fitted G-computation gives -15.3 points, and the doubly robust AIPW estimator gives -16.2 points with influence-function standard error 3.9 (95% CI -23.9 to -8.5). The consistency across a design-free estimator, an outcome-model-based estimator, and a doubly robust estimator is exactly what one hopes to see in a randomized comparison, and the magnitudes match the intention-to-treat analysis of Section 4.1.

Conditional effects: where the benefit lives

The cross-fitted T-learner decomposes this average into patient-level predicted benefits (Figure 18). The distribution is wide: predicted five-year risk reductions range from negligible to

Table 10: Average treatment effect on the five-year DFS event risk (negative = risk reduction with levamisole plus 5-FU).

Estimator	ATE (percentage points)	SE	95% CI
Kaplan–Meier difference (bootstrap)	−16.8	4.0	−24.5 to −9.1
G-computation (cross-fitted, pseudo-obs)	−15.3	0.5	−16.2 to −14.4
AIPW (cross-fitted, doubly robust)	−16.2	3.9	−23.9 to −8.5

over 40 percentage points, with quartile means of 28.9, 18.6, 12.6, and 1.1 points. Crucially, the predicted benefit is strongly tied to baseline risk (correlation 0.648 between control-arm predicted risk and predicted benefit): the patients predicted to benefit most are those whose disease is most aggressive—high nodal burden, advanced local extent—while low-risk patients have little room for absolute improvement even under a proportionally constant treatment effect. Table 11 sharpens the picture with observed outcomes: within the highest predicted-benefit quartile (mean baseline risk 71.0%), the observed five-year risk difference is −12.8 points; within the lowest quartile (baseline risk 46.6%), it is −19.3 points—statistically compatible with the same relative effect everywhere. The apparent paradox (large predicted but modest observed benefit in the top quartile, the reverse in the bottom) reflects both estimation noise in quartile-specific contrasts (each based on ~64–98 events) and the mechanical relationship between baseline risk and achievable absolute benefit.




Figure 18: Predicted absolute benefit of levamisole plus 5-FU at five years. Left: distribution across patients (AIPW average marked). Right: predicted benefit versus predicted baseline risk, coloured by randomized arm.

Table 11: Predicted-benefit quartiles from the cross-fitted T-learner. “Predicted benefit” is the model-implied five-year risk reduction; “observed risk difference” is the corresponding pseudo-observation contrast by arm within quartile.

Quartile	<i>n</i>	Pred. benefit (pts)	Baseline risk	Observed diff (pts)	Events
Q1 (largest predicted benefit)	149	28.9	71.0%	−12.8	98
Q2	148	18.6	58.2%	−16.1	79
Q3	148	12.6	52.0%	−11.8	72
Q4 (smallest)	149	1.1	46.6%	−19.3	64

The synthesis is the section’s key message: in this trial, heterogeneity in *absolute* benefit is real and large, but it is driven almost entirely by heterogeneity in *baseline risk* rather than by

qualitative effect modification—consistent with the null interaction tests of Section 4.2. For clinical decision-making, this is the most useful kind of heterogeneity: it means a single relative effect, applied to an accurate prognostic model, suffices to identify who gains most in absolute terms, with no need to invoke distinct treatment-response biology. All such CATE statements remain exploratory in a single trial of this size.

5 Discussion

5.1 Principal findings

Six analytical components, aimed at one 1980s dataset, converge on one conclusion and enrich it from different angles. The conclusion: levamisole plus 5-FU after resection of stage B/C colon carcinoma delivers a large, robust treatment benefit on disease-free survival—hazard ratio 0.62 adjusted, 0.59–0.61 Bayesian, 0.60 under competing risks—corresponding to a five-year absolute gain of 16.8 percentage points, a number needed to treat of six, and 0.63 disease-free years gained in restricted mean survival time. The enrichment: the effect is stable across seven pre-specified subgroups with no credible interaction; a modern group-sequential design would likely have stopped the trial for efficacy at its first interim analysis (HR 0.52, $Z = 3.38$ against a boundary of 2.96), and Bayesian monitoring would have signalled benefit at a quarter of the evidence; covariate adjustment would have bought real power at no inferential cost; and the largest absolute benefits accrue to the highest-risk patients through the arithmetic of baseline risk rather than differential response.

The re-enacted monitoring result deserves particular emphasis because it connects historical data to contemporary trial ethics. Fifty percent of the 506 DFS events in this dataset contain essentially all the evidence needed for the superiority conclusion; the remaining half of the patients’ events—years of follow-up and hundreds of recurrences—added precision but not decisions. Translated to the 1980s context in which the 625 patients randomized to observation or levamisole alone in this very trial continued on an ineffective strategy, an early stop would have redirected patients to the active regimen years sooner. The counterfactual is stylized (the historical investigators lacked today’s data-management speed and regulatory framing), but it illustrates with unusual concreteness why alpha-spending and beta-spending designs became the default.

5.2 Methodological lessons

Several results function as transferable teaching points. **Non-collapsibility**, usually a paragraph in a textbook, appears here as a measured quantity: with trial-calibrated prognostic strength, the unadjusted Cox estimator targets a marginal hazard ratio attenuated by 0.033 on the log scale, so unadjusted and adjusted estimates are *different estimands* even in a perfectly randomized trial (Simulation Study A). The RMST analysis supplies the constructive counterpoint: because the restricted-mean difference is collapsible, its adjusted and unadjusted estimates coincide (0.57 vs. 0.63 years within a five-year horizon, well within sampling error), making RMST an attractive companion estimand whenever hazard-ratio interpretations become delicate (Royston and Parmar, 2011; Uno et al., 2014; Martinussen and Vansteelandt, 2013). **Covariate adjustment** raises power from 72.6% to 80.9% in exactly the scenario regulators care about—modest effects, realistic sample sizes—which is the quantitative argument behind the FDA/ICH position that adjustment should be pre-specified and reported even when unadjusted results match (U.S. Food and Drug Administration, 2023; Kahan et al., 2014). **Competing risks** change the estimand, not necessarily the conclusion: the naive one-minus-Kaplan–Meier overstates cumulative recurrence by about one percentage point at eight years, yet the Fine–Gray and cause-specific readings agree on the treatment effect. **Honest**

validation matters more than model sophistication: the apparent-to-cross-validated optimism gap reaches 0.12 in the C -index for the most aggressive learner, while a clinical Cox model loses nothing—a reminder that in small-to-moderate clinical datasets, disciplined regression frequently matches ensembles. **Bayesian monitoring** demonstrates both its appeal (decisions at 25% of the information) and its governance question (posterior probabilities do not self-correct for repeated looks, so their use in confirmatory settings requires either calibration to frequentist error rates or a genuinely decision-theoretic framing).

5.3 Strengths

The project’s strengths are its integration and its reproducibility. Every module runs from version-controlled code with fixed seeds, communicates exclusively through files, and regenerates every number and figure in this paper via a single `make` invocation; all 34 tabular outputs of the original pipeline were re-derived during manuscript preparation and verified to be byte-identical to the committed versions. Quality-control assertions are embedded in the data pipeline (record counts, within-patient covariate consistency, endpoint logic), analyses are pre-specified in structure (primary composite endpoint, adjustment set, subgroup list), and sensitivity analyses were run in both directions—towards stronger assumptions (parametric Weibull Bayesian model) and weaker ones (competing risks, splines, complete case versus imputation, RMST). The simulation studies were calibrated to the real data-generating mechanism rather than to stylized parameters, so their conclusions transfer directly to trials of this type.

5.4 Limitations

Five limitations bound the interpretation. First, this is a *re-analysis of a single historical trial* from an era of different supportive care and staging; the quantitative findings (especially absolute risks) should not be read as contemporary practice, and levamisole itself was later displaced by superior regimens. Second, the subgroup and CATE analyses are exploratory: with only 324 events in the two-arm comparison, interaction tests have limited power, the T-learner’s quartile contrasts rest on 60–100 events each, and no external validation cohort was available to check the predicted benefit ranking; the honest claim is methodological demonstration, not clinical deployment. Third, the Weibull baseline is a parametric choice; although posterior predictive checks pass and LOO prefers it over the exponential, a piecewise-exponential or semi-parametric Bayesian baseline would be a natural robustness extension, as would a time-varying treatment effect in the Bayesian model given the covariate-level PH violations detected. Fourth, the re-enacted monitoring and Bayesian sequential analyses re-use the final data to reconstruct interim states; the reconstruction censors follow-up at a common time aligned to the event count and is faithful to the ordering of the event stream, but the exercise remains a counterfactual simulation, and early-stopping estimates from such re-enactments are biased toward larger effects (as the interim HR of 0.52 versus the final 0.62 illustrates). Fifth, the prediction models use only ten baseline covariates from the original case-report forms; modern depth of data (genomics, imaging) would shift the accuracy ceiling, though the overfitting lesson of the apparent-versus-validated contrast would only sharpen.

A further, more technical caveat concerns the simulated operating characteristics of the sequential design (Study B): the simulated rejection probability under the alternative (82.9%) falls short of the design-theoretic 90% because the simulation couples an unadjusted primary analysis (whose marginal hazard ratio is attenuated by the strong calibration covariates, per Study A) with a simulation sample size calibrated to the conditional effect, and because futility adherence removes some would-be winners. Designs that pair covariate-adjusted primary analyses with group-sequential monitoring would recover part of this gap—a combination now explicitly

encouraged by regulators (U.S. Food and Drug Administration, 2023) and worth studying in its own right.

5.5 Conclusions

This re-analysis set out to demonstrate what a modern biostatistical toolkit, applied with discipline to a classic randomized trial, can extract beyond the original publication—and the answer is: the same conclusion, reached faster (sequential design re-enactment), quantified more honestly (imputation, competing risks, absolute effects in both percentage points and disease-free years), understood more deeply (simulation-calibrated operating characteristics, Bayesian robustness, collapsibility made concrete), and personalized in direction (risk-driven absolute benefit). The accompanying repository delivers the full pipeline—R and Python, 23 figures, 37 tabular outputs, and this paper—reproducible with one command, in the hope that it serves both as a reference implementation and as a teaching resource on how frequentist, Bayesian, and machine-learning perspectives can be made to interrogate the same data in complementary ways.

Data availability

The trial data are the `colon` dataset distributed with the R package `survival` (Therneau, 2026), derived from the NCCTG adjuvant colon cancer program (Laurie et al., 1989; Moertel et al., 1990) and obtained from the Rdatasets mirror; they are redistributed with the repository in `data/raw/colon.csv` with full attribution for research and educational use.

Code availability

All analysis code (six R modules and four Python modules), the reproducibility environment files, and every figure and table in this paper are generated from the accompanying repository (`Makefile: make all`), with fixed random seeds (global seed 20260916) and module-to-module communication exclusively through versioned files in `data/` and `outputs/`. Exact software versions are recorded in `environment/r_session_info.txt` and `environment/python_package_versions.txt`.

References

- Odd O. Aalen. Nonparametric inference for a family of counting processes. *Annals of Statistics*, 6(4):701–726, 1978. doi:10.1214/aos/1176344247.
- Per K. Andersen, Jan P. Klein, and Susanne Rosthøj. Generalised linear models for correlated pseudo-observations, with applications to multi-state models. *Biometrika*, 90(1):15–27, 2003. doi:10.1093/biomet/90.1.15.
- Keaven M. Anderson. *gsDesign: Group Sequential Design*, 2026. URL <https://CRAN.R-project.org/package=gsDesign>. R package version 3.11.0. CRAN.R-project.org/package=gsDesign.
- Jason P. Fine and Robert J. Gray. A proportional hazards model for the subdistribution of a competing risk. *Journal of the American Statistical Association*, 94(446):496–509, 1999. doi:10.1080/01621459.1999.10474144.
- Andrew Gelman, John B. Carlin, Hal S. Stern, David B. Dunson, Aki Vehtari, and Donald B. Rubin. *Bayesian Data Analysis*. CRC Press, Boca Raton, 3rd edition, 2013. doi:10.1201/b16618.

- Erika Graf, Claudia Schmoor, Willi Sauerbrei, and Martin Schumacher. Assessment and comparison of prognostic classification schemes for survival data. *Statistics in Medicine*, 18 (17–18):2529–2545, 1999. pubmed.ncbi.nlm.nih.gov/10474150/.
- Patricia M. Grambsch and Terry M. Therneau. Proportional hazards tests and diagnostics based on weighted residuals. *Biometrika*, 81(3):515–526, 1994. [doi:10.1093/biomet/81.3.515](https://doi.org/10.1093/biomet/81.3.515).
- Jr. Harrell, Frank E. *Regression Modeling Strategies: With Applications to Linear Models, Logistic and Ordinal Regression, and Survival Models*. Springer, Cham, 2nd edition, 2015. [doi:10.1007/978-3-319-19425-7](https://doi.org/10.1007/978-3-319-19425-7).
- Miguel A. Hernán and James M. Robins. *Causal Inference: What If*. Chapman & Hall/CRC, Boca Raton, 2020. www.hsph.harvard.edu/miguel-hernan/causal-inference-book/.
- Irving K. Hwang, Weichung J. Shih, and John S. De Cani. Group sequential designs using a family of type I error probability spending functions. *Statistics in Medicine*, 9(12):1439–1445, 1990. [doi:10.1002/sim.4780091210](https://doi.org/10.1002/sim.4780091210).
- Brennan C. Kahan, Vipul Jairath, Caroline J. Doré, and Tim P. Morris. The risks and rewards of covariate adjustment in randomized trials: an assessment of 12 outcomes from 8 studies. *Trials*, 15:139, 2014. [doi:10.1186/1745-6215-15-139](https://doi.org/10.1186/1745-6215-15-139).
- K. K. Gordon Lan and David L. DeMets. Discrete sequential boundaries for clinical trials. *Biometrika*, 70(3):659–663, 1983. [doi:10.1093/biomet/70.3.659](https://doi.org/10.1093/biomet/70.3.659).
- J. A. Laurie, C. G. Moertel, T. R. Fleming, H. S. Wieand, J. E. Leigh, J. Rubin, P. J. Novotny, L. K. Everson, J. E. Krook, J. Malliard, J. G. Gerstner, and J. A. Mailliard. Surgical adjuvant therapy of large-bowel carcinoma: an evaluation of levamisole and the combination of levamisole and fluorouracil. *Journal of Clinical Oncology*, 7(10):1447–1456, 1989. [doi:10.1200/JCO.1989.7.10.1447](https://doi.org/10.1200/JCO.1989.7.10.1447).
- Torben Martinussen and Stijn Vansteelandt. On collapsibility and confounding bias in Cox and Aalen regression models. *Lifetime Data Analysis*, 19(3):279–296, 2013. [doi:10.1007/s10985-013-9242-z](https://doi.org/10.1007/s10985-013-9242-z).
- C. G. Moertel, T. R. Fleming, J. S. Macdonald, D. G. Haller, J. A. Laurie, P. J. Goodman, J. S. Ungerleider, W. A. Emerson, D. C. Tormey, J. H. Glick, M. H. Veeder, and J. A. Mailliard. Levamisole and fluorouracil for adjuvant therapy of resected colon carcinoma. *New England Journal of Medicine*, 322(6):352–358, 1990. [doi:10.1056/NEJM199002083220606](https://doi.org/10.1056/NEJM199002083220606).
- Peter C. O’Brien and Thomas R. Fleming. A multiple testing procedure for clinical trials. *Biometrics*, 35(3):549–556, 1979. [doi:10.2307/2530245](https://doi.org/10.2307/2530245).
- Esben T. Parner and Per K. Andersen. Regression analysis of censored data using pseudo-observations. *The Stata Journal*, 10(3):408–422, 2010. [doi:10.1177/1536867X1001000310](https://doi.org/10.1177/1536867X1001000310).
- Patrick Royston and Mahesh K. B. Parmar. The use of restricted mean survival time to estimate the treatment effect in randomized clinical trials when the proportional hazards assumption is in doubt. *Statistics in Medicine*, 30(20):2409–2421, 2011. [doi:10.1002/sim.4274](https://doi.org/10.1002/sim.4274).
- Donald B. Rubin. *Multiple Imputation for Nonresponse in Surveys*. Wiley, New York, 1987. [doi:10.1002/9780470316696](https://doi.org/10.1002/9780470316696).
- John Salvatier, Thomas V. Wiecki, and Christopher Fonnesbeck. Probabilistic programming in Python using PyMC3. *PeerJ Computer Science*, 2:e55, 2016. [doi:10.7717/peerj-cs.55](https://doi.org/10.7717/peerj-cs.55); PyMC 6: <https://www.pymc.io>.

- David A. Schoenfeld. Sample-size formula for the proportional-hazards regression model. *Biometrics*, 39(2):499–503, 1983. doi:10.2307/2531101.
- Jonathan A. C. Sterne, Ian R. White, John B. Carlin, Michael Spratt, Patrick Royston, Michael G. Kenward, Angela M. Wood, and James R. Carpenter. Multiple imputation for missing data in epidemiological and clinical research: potential and pitfalls. *BMJ*, 338: b2393, 2009. doi:10.1136/bmj.b2393.
- Terry M. Therneau. *A Package for Survival Analysis in R*, 2026. URL <https://CRAN.R-project.org/package=survival>. R package version 3.8-12. CRAN.R-project.org/package=survival.
- Hajime Uno, Brian Claggett, Lu Tian, Eisuke Inoue, Paul Gallo, Takashi Miyata, Deborah Schrag, Masahiro Takeuchi, Yuki Uyama, Lihui Zhao, Hicham Skali, Scott Solomon, Scott Jacobus, Michael Hughes, Milton Packer, and Lee-Jen Wei. Moving beyond the hazard ratio in quantifying the between-group difference in survival analysis. *Journal of Clinical Oncology*, 32(22):2380–2385, 2014. doi:10.1200/jco.2014.55.2208.
- U.S. Food and Drug Administration. Adjusting for covariates in randomized clinical trials for drugs and biologics with continuous outcomes. Guidance for industry, U.S. Food and Drug Administration, 2023. www.fda.gov/regulatory-information/search-fda-guidance-documents/adjusting-covariates-randomized-clinical-trials-drugs-and-biologics-continuous-outcomes.
- Stef van Buuren and Karin Groothuis-Oudshoorn. mice: Multivariate imputation by chained equations in R. *Journal of Statistical Software*, 45(3):1–67, 2011. doi:10.18637/jss.v045.i03.
- Aki Vehtari, Andrew Gelman, and Jonah Gabry. Practical Bayesian model evaluation using leave-one-out cross-validation and WAIC. *Statistics and Computing*, 27(5):1413–1432, 2017. doi:10.1007/s11222-016-9696-4.
- Ian R. White and Patrick Royston. Imputing missing covariate values for the Cox model. *Statistics in Medicine*, 28(15):1982–1998, 2009. doi:10.1002/sim.3618.